


Transfer to natural cgreGFP after training on amacGFP + ppluGFP



Each dot is one held-out natural cgreGFP sequence ($n = 4,904$); the dashed line is perfect prediction.

Training: 7,355 amacGFP + 7,354 ppluGFP; validation: 2,452 + 2,451; test: 4,904 cgreGFP. No labeled cgreGFP was used for training or checkpoint selection.

Aubin/MLPs: aligned one-hot inputs. CNN Jannis: full residue-level ESM-2 embeddings. Rows are independent fits and are not averaged.