













Natural cgreGFP → natural cgreGFP

Aubin 1-10-1

Small MLP

Deep MLP

CNN Jannis



Measured log10 fluorescence

Measured log10 fluorescence

Measured log10 fluorescence

Measured log10 fluorescence

Each dot is one held-out natural cgreGFP sequence ($n = 4,904$); the dashed line is perfect prediction.

Training / validation / test: 14,709 / 4,903 / 4,904 $\approx$ 60% / 20% / 20%. The split and test sequences are identical across rows.

Aubin/MLPs: aligned one-hot inputs. CNN Jannis: full residue-level ESM-2 embeddings. Rows are independent fits and are not averaged.