



Natural cgreGFP: measured vs predicted · seed 43


Aubin 1-10-1




Small MLP



Deep MLP



CNN Jannis



Each dot: one held-out sequence ($n = 4,904$); dashed line: perfect prediction.

Train / validation / test: 14,709 / 4,903 / 4,904 $\approx$ 60% / 20% / 20%.

cgre-only transfer controls: aligned one-hot Aubin/MLPs; full native ESM embeddings for CNN.

Individual run predictions; seeds are shown separately, not averaged.